

Santa Clara University
COEN 241: Cloud Computing
HW 1: System vs OS Virtualization
Ankita Kumar (W1620637)

HOST ENVIRONMENT CONFIGURATION:

Detailed system configurations (CPU, Mem, etc...) on which all the experiments have been performed.



Hardware Overview:

Model Name:	MacBook Air
Model Identifier:	MacBookAir10,1
Chip:	Apple M1
Total Number of Cores:	8 (4 performance and 4 efficiency)
Memory:	8 GB
System Firmware Version:	6723.101.4
OS Loader Version:	6723.101.4
Activation Lock Status:	Enabled

QEMU INSTALLATION ON M1 & IT'S CONFIGURATION:

Below are the main steps to install QEMU manually:

- a. Install Xcode command line tools.

```
xcode-select --install
```

- b. Install Homebrew arm64

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"  
echo 'eval "$(/opt/homebrew/bin/brew shellenv)"' >> /Users/${(logname)}/.zprofile  
eval "$(/opt/homebrew/bin/brew shellenv)"
```

- c. Install necessary packages for building

```
brew install libffi gettext glib pkg-config autoconf automake pixman ninja
```

- d. Clone Qemu

```
git clone https://github.com/qemu/qemu
```

- e. Change directory to Qemu repository

```
cd qemu
```

- f. Checkout to commit dated June 03, 2021 v6.0.0

```
git checkout 3c93dfa42c394fdd55684f2fbf24cf2f39b97d47
```

- g. Apply patch series v8 by Alexander Graf

```
curl https://patchwork.kernel.org/series/485309/mbox/ | git am
```

h. Building Qemu installer

```
mkdir build && cd build  
../configure --target-list=aarch64-softmmu  
make -j8
```

i. Install Qemu

```
sudo make install
```

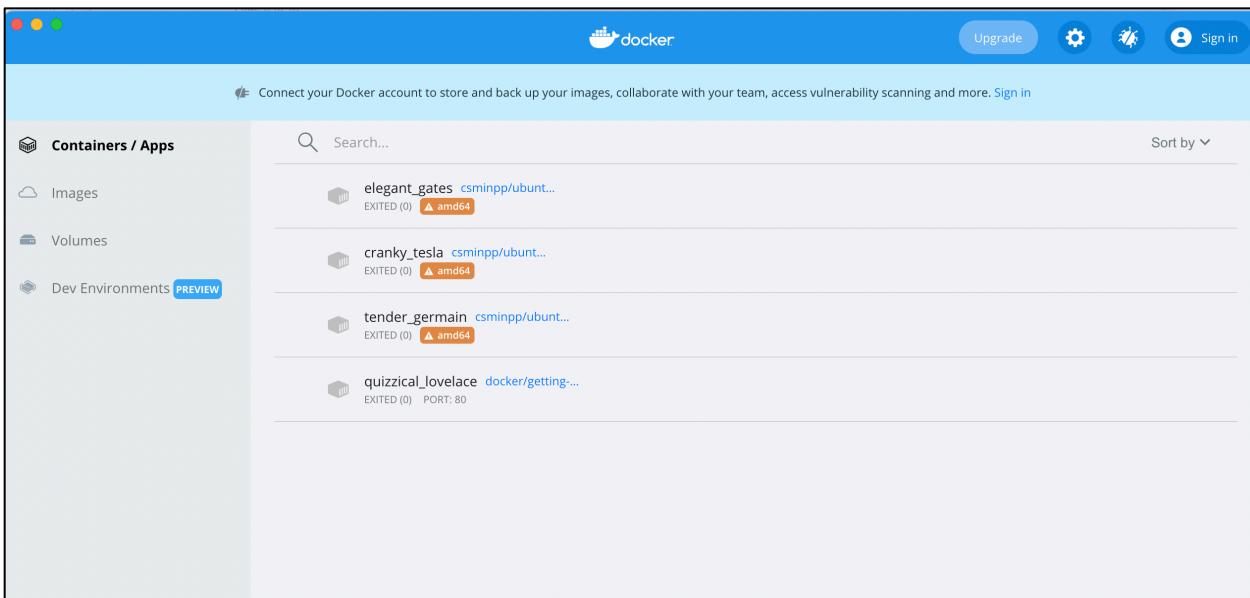
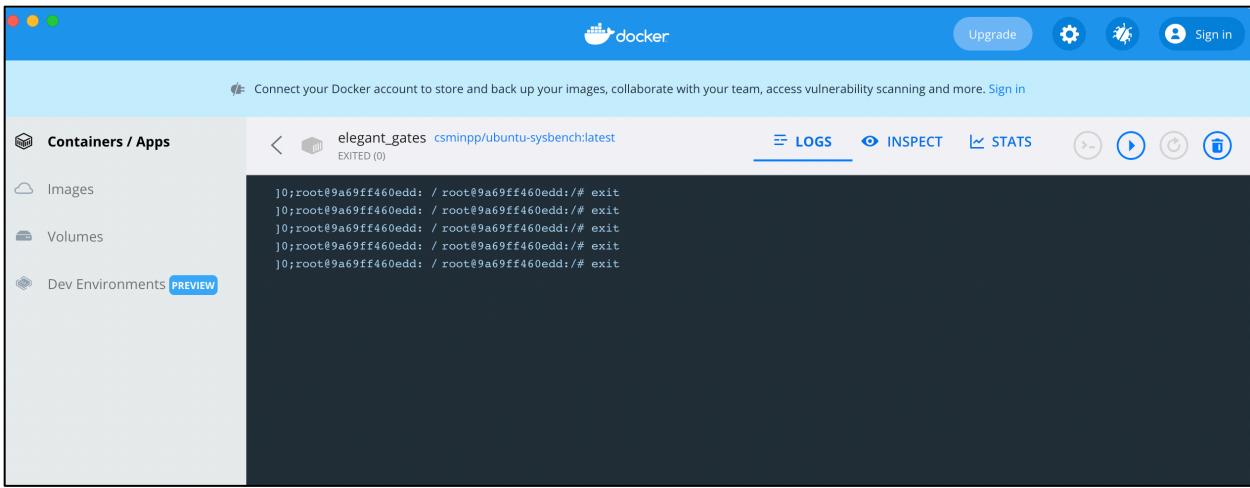
DOCKER INSTALLATION ON M1 & IT'S CONFIGURATION:

Installing Docker was straight forward through this link.

[Docker Desktop for Apple silicon | Docker Documentation](https://docs.docker.com/desktop/mac/apple-silicon/)

Further directly using the CLI to execute the scripts.

The screenshot shows a web browser displaying the Docker Desktop for Apple silicon documentation on the docker docs website. The URL in the address bar is <https://docs.docker.com/desktop/mac/apple-silicon/>. The page has a blue header with navigation links for Home, Guides, Manuals, Reference, and Samples. The 'Manuals' link is currently selected. Below the header, there's a breadcrumb trail: Manuals / Docker Desktop / Mac / Apple silicon. The main content area is titled 'Docker Desktop for Apple silicon' and includes a sub-section 'Apple silicon'. It features a sidebar with links like 'Logs and troubleshooting', 'Release notes', 'Previous versions', 'Windows', 'Dashboard', 'Dev Environments (Preview)', 'Multi-arch support', 'Deploy on Kubernetes', 'FAQs', 'Backup and restore data', 'Container Engine', 'Container Compose', and 'Container Hub'. A note states 'Estimated reading time: 3 minutes'. The main text explains that Docker Desktop for Mac on Apple silicon is now available as a GA release, enabling development with local environments and ARM-based applications. It also mentions multi-platform images and the Docker Hub. A large blue button at the bottom right says 'Download Docker Desktop' with the sub-option 'Mac with Apple chip' selected. To the right of the main content, there are sections for 'Edit this page', 'Request docs changes', and 'On this page' with links to 'System requirements', 'Known issues', and various fix history sections for RC 3, RC 2, RC 1, and preview 3.1.0, along with a 'Feedback' link.



QEMU VM ON M1 & IT'S CONFIGURATION:

Create Ubuntu Server using QEMU on Silicon based Apple Macs. First run, close manually after installation aka first reboot.

1. Install ISO library using below CLI:

```
curl https://cdimage.ubuntu.com/releases/20.04/release/ubuntu-20.04.2-live-server-arm64.iso -o ubuntu-lts.iso
```

2. Create Virtual Hardisk on VM of size 25 GB.

```
qemu-img create -f qcow2 virtual-disk.qcow2 25G
```

3. Creating File for UEFI variables.

```
dd if=/dev/zero conv=sync bs=1m count=64 of=ovmf_vars.fd
```

4. Run the below CLI argument to install VM with these configs.

```
qemu-system-aarch64 \
    -machine virt,accel=hvf,highmem=off \
    -cpu cortex-a72 -smp 4 -m 4G \
    -device virtio-gpu-pci \
    -device virtio-keyboard-pci \
    -drive "format=raw,file=edk2-aarch64-code.fd,if=pflash,readonly=on" \
    -drive "format=raw,file=ovmf_vars.fd,if=pflash" \
    -drive "format=qcow2,file=virtual-disk.qcow2" \
    -cdrom ubuntu-lts.iso
```

5. The above commands will download Ubuntu Image file. Complete installation manually. Run the above same command line arguments again without -cdrom parameter.

Second run for ubuntu server:

```
qemu-system-aarch64 \
-machine virt,accel=hvf,highmem=off \
-cpu cortex-a72 -smp 4 -m 4G \
-device virtio-gpu-pci \
-device virtio-keyboard-pci \
-drive "format=raw,file=edk2-aarch64-code.fd,if=pflash,readonly=on" \
-drive "format=raw,file=ovmf_vars.fd,if=pflash" \
-drive "format=qcow2,file=virtual-disk.qcow2"
```

The screenshot shows a terminal window titled "QEMU - (Press ctrl + alt + g to release Mouse)". The window displays the following text:

```
Ubuntu 20.04.3 LTS akumar tty1
akumar login: akumar
Password:
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.4.0-89-generic aarch64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

System information as of Fri 22 Oct 2021 09:59:43 PM UTC

System load:          1.1
Usage of /:            31.5% of 19.56GB
Memory usage:          5%
Swap usage:            0%
Processes:             146
Users logged in:      0
IPv4 address for enp0s1: 10.0.2.15
IPv6 address for enp0s1: fec0::5054:ff:fe12:3456

27 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Last login: Fri Oct 22 15:09:50 UTC 2021 from 10.0.2.2 on pts/0
akumar@akumar:~$ -
```

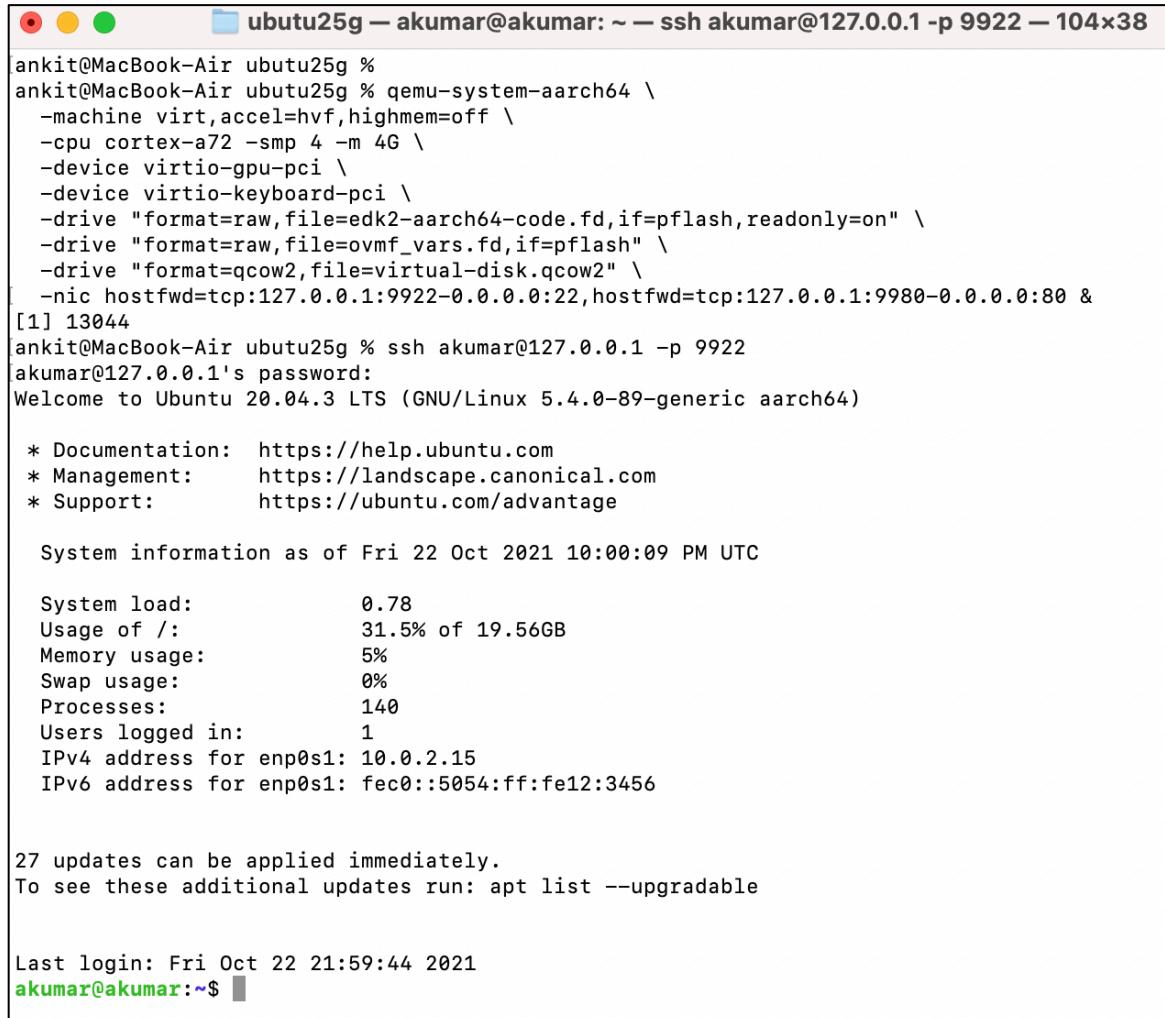
CONNECT UBUNTU VIA SSH AND SERVE PORT 80 AT LOCALHOST.

1. Start using server:

```
qemu-system-aarch64 \
-machine virt,accel=hvf,highmem=off \
-cpu cortex-a72 -smp 4 -m 4G \
-device virtio-gpu-pci \
-device virtio-keyboard-pci \
-drive "format=raw,file=edk2-aarch64-code.fd,if=pflash,readonly=on" \
-drive "format=raw,file=ovmf_vars.fd,if=pflash" \
-drive "format=qcow2,file=virtual-disk.qcow2" \
-nic hostfwd=tcp:127.0.0.1:9922-0.0.0.0:22,hostfwd=tcp:127.0.0.1:9980-0.0.0.0:80 &
```

2. Login via SSH using –

```
ssh akumar@127.0.0.1 -p 9922
```



The screenshot shows a terminal window titled "ubuntu25g — akumar@akumar: ~ — ssh akumar@127.0.0.1 -p 9922 — 104x38". The session starts with a command to start a QEMU system, followed by the password prompt and a successful login. The user then runs "systemctl status" to check system information, which includes load average, memory usage, swap usage, processes, and network interfaces. It also shows 27 upgradeable packages and the last login time. The terminal ends with a prompt at the bottom.

```
ankit@MacBook-Air ubutu25g %
ankit@MacBook-Air ubutu25g % qemu-system-aarch64 \
-machine virt,accel=hvf,highmem=off \
-cpu cortex-a72 -smp 4 -m 4G \
-device virtio-gpu-pci \
-device virtio-keyboard-pci \
-drive "format=raw,file=edk2-aarch64-code.fd,if=pflash,readonly=on" \
-drive "format=raw,file=ovmf_vars.fd,if=pflash" \
-drive "format=qcow2,file=virtual-disk.qcow2" \
-nic hostfwd=tcp:127.0.0.1:9922-0.0.0:22,hostfwd=tcp:127.0.0.1:9980-0.0.0:80 &
[1] 13044
ankit@MacBook-Air ubutu25g % ssh akumar@127.0.0.1 -p 9922
akumar@127.0.0.1's password:
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.4.0-89-generic aarch64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

 System information as of Fri 22 Oct 2021 10:00:09 PM UTC

 System load:          0.78
 Usage of /:           31.5% of 19.56GB
 Memory usage:         5%
 Swap usage:           0%
 Processes:            140
 Users logged in:     1
 IPv4 address for enp0s1: 10.0.2.15
 IPv6 address for enp0s1: fec0::5054:ff:fe12:3456

27 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Last login: Fri Oct 22 21:59:44 2021
akumar@akumar:~$
```

3. Now we can directly run execute the shell scripts at local host by downloading Sysbench using below argument:

```
$ sudo apt-get install sysbench
```

Test different arguments of QEMU :

1. Emulated Machine

-M <machine-type>
-machine <machine-type>[,<property>[=<value>][,...]]

Machine

- `-machine type=q35,accel=kvm` - Modern chipset (PCIe, AHCI, ...) and hardware virtualization acceleration
- `-object rng-random,id=rng0,filename=/dev/urandom -device virtio-rng-pci,rng=rng0` - Pass-through for host random number generator. Accelerates startup of e.g. Debian VMs because of missing entropy.

2. Processor Type

-cpu <model>[,<FEATURE>][...]

3. Processor Topology

smp <n>[,cores=<ncores>][,threads=<nthreads>][,sockets=<nsocks>][,maxcpus=<maxcpus>]

Hypervisor and guest operating system limits on processor topology apply.

Processor

- `-cpu <CPU>` - Specify a processor architecture to emulate. To see a list of supported architectures, run: `qemu-system-x86_64 -cpu ?`
- `-cpu host` - (Recommended) Emulate the host processor.
- `-smp <NUMBER>` - Specify the number of cores the guest is permitted to use. The number can be higher than the available cores on the host system. Use `-smp $(nproc)` to use all currently available cores.

4. Memory Size

-m <megs>

RAM

- `-m MEMORY` - Specify the amount of memory (default: 128 MB). For instance: `-m 256M` (*M* stands for Megabyte, *G* for Gigabyte).

5. Generic Drive:

-drive <option>[,<option>[,<option>[,...]]]

Hard drive

- `-hda IMAGE.img` - Set a virtual hard drive and use the specified image file for it.
- `-drive` - Advanced configuration of a virtual hard drive:

6. Optical Drive:

- `-cdrom IMAGE.iso` - Set a virtual CDROM drive and use the specified image file for it.
- `-cdrom /dev/cdrom` - Set a virtual CDROM drive and use the host drive for it.
- `-drive` - Advanced configuration of a virtual CDROM drive:

CPU PERFORMANCE TESTING:

- We are testing CPU performance using sysbench tool in both Ubuntu VM and Docker Container for below CPU parameters: **cpu-max-prime=(10000 .. 80000)th. & THREAD =1**
- Each Sysbench measurement is repeated at least 5 times for each test case

Automation for FILE IO performance test

Automation for CPU performance test is done using Shell script program.

```
Executable File | 18 lines (17 sloc) | 325 Bytes

1  #!/bin/bash
2  #Loop
3  exec 3>&1 4>&2
4  trap 'exec 2>&4 1>&3' 0 1 2 3
5  exec 1>Docker_log.output 2>&1
6  START=1
7  END=5
8  cpm_arr=(10000 20000 30000 40000 50000 60000 70000 80000)
9  for i in "${cpm_arr[@]}"
10 do
11 echo "CMP Value:" $i
12 for j in $(eval echo "{$START..$END}")
13 do
14 echo "Loop" $j
15 sysbench --test=cpu --cpu-max-prime=$i run
16 done
17 done
18
```

TEST CASE 1:

- For a single thread

sysbench --test=cpu --cpu-max-prime=10000 run

```
ankit@MacBook-Air COEN241-HW1 % cat Docker_log_CPU.output
CMP Value: 10000
Loop 1
WARNING: the --test option is deprecated. You can pass a
sysbench 1.0.20 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 10000

Initializing worker threads...
Threads started!

CPU speed:
events per second: 1000.41

General statistics:
total time: 10.00014s
total number of events: 10001

Latency (ms):
min: 0.46
avg: 0.99
max: 2.55
95th percentile: 1.04
sum: 9938.85

Threads fairness:
events (avg/stddev): 10014.0000/0.00
execution time (avg/stddev): 9.9389/0.00

Loop 2
WARNING: the --test option is deprecated. You can pass a script
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time
```

```
Prime numbers limit: 10000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 1000.41

General statistics:
total time: 10.0014s
total number of events: 10014

Latency (ms):
min: 0.46
avg: 0.99
max: 2.55
95th percentile: 1.04
sum: 9938.85

Threads fairness:
events (avg/stddev): 10014.0000/0.00
execution time (avg/stddev): 9.9389/0.00

Loop 2
WARNING: the --test option is deprecated. You can pass a script
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time
```

Docker container

QEMU VM

For QEMU

Test run	Prime numbers limit	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	10000	10.0000s	0	0	0.06	13,672,861.69
2	10000	10.0000s	0	0	0.04	13,695,972.75
3	10000	10.0000s	0	0	0.06	13,744,287.78
4	10000	10.0000s	0	0	0.05	13,743,972.53
5	10000	10.0001s	0	0	0.05	13,741,889.82

For Docker

A	B	C	D	E	F	G
Test run	CMP	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	10000	10.0014	0.46	0.99	2.55	1000.41
2	10000	10.001	0.46	1	5.48	994.79
3	10000	10.001	0.46	0.98	3.67	1018.5
4	10000	10.001	0.46	1	4.06	997.51
5	10000	10.0005	0.46	1.01	25.76	987.69

TEST CASE 2 :

- For a single thread

sysbench --test=cpu --cpu-max-prime=20000 run

```
CMP Value: 20000
Loop 1
WARNING: the --test option is deprecated. You can pass
sysbench 1.0.20 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 13741889.82

General statistics:
total time: 10.0001s
total number of events: 137424821

Latency (ms):
min: 0.00
avg: 0.00
max: 0.05
95th percentile: 0.00
sum: 3775.06

Threads fairness:
events (avg/stddev): 137424821.0000/0.00
execution time (avg/stddev): 3.7751/0.00
```

```
CMP Value: 20000
Loop 1
WARNING: the --test option is deprecated. You can pass a s
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 388.38

General statistics:
total time: 10.0020s
total number of events: 3886

Latency (ms):
min: 1.19
avg: 2.56
max: 5.32
95th percentile: 2.66
sum: 9963.47

Threads fairness:
events (avg/stddev): 3886.0000/0.00
execution time (avg/stddev): 9.9635/0.00
```

Docker container

QEMU VM

For Docker

A	B	C	D	E	F	G
Test run	Prime numbers limit	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	20000	10.0001s	0	0	0.05	13,741,889.82
2	20000	10.0001s	0	0	0.05	13,736,544.34
3	20000	10.0001s	0	0	0.06	13,744,655.94
4	20000	10.0001s	0	0	0.05	13,742,488.02
5	20000	10.0001s	0	0	0.04	13,572,744.02

For QEMU

Test run	CMP	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	20000	10.002	1.19	2.56	5.32	388.38
2	20000	10.0008	1.19	2.53	8.43	393.93
3	20000	10.0028	1.19	2.57	4.96	386.65
4	20000	10.002	1.19	2.57	4.06	387.64
5	20000	10.0008	1.19	2.56	8.04	388.34

TEST CASE 3 :

- **For a single thread**

sysbench --test=cpu --cpu-max-prime=80000 run

```
CMP Value: 80000
Loop 1
WARNING: the --test option is deprecated. You can pass a
sysbench 1.0.20 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 80000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 13413099.56

General statistics:
total time: 10.0001s
total number of events: 134136495

Latency (ms):
min: 0.00
avg: 0.00
max: 0.06
95th percentile: 0.00
sum: 3794.00

Threads fairness:
events (avg/stddev): 134136495.0000/0.00
execution time (avg/stddev): 3.7940/0.00
```

```
CMP Value: 80000
Loop 1
WARNING: the --test option is deprecated. You can pass a
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 80000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 56.93

General statistics:
total time: 10.0078s
total number of events: 570

Latency (ms):
min: 11.96
avg: 17.54
max: 22.76
95th percentile: 17.95
sum: 9997.29

Threads fairness:
events (avg/stddev): 570.0000/0.00
execution time (avg/stddev): 9.9973/0.00
```

Docker container

QEMU VM

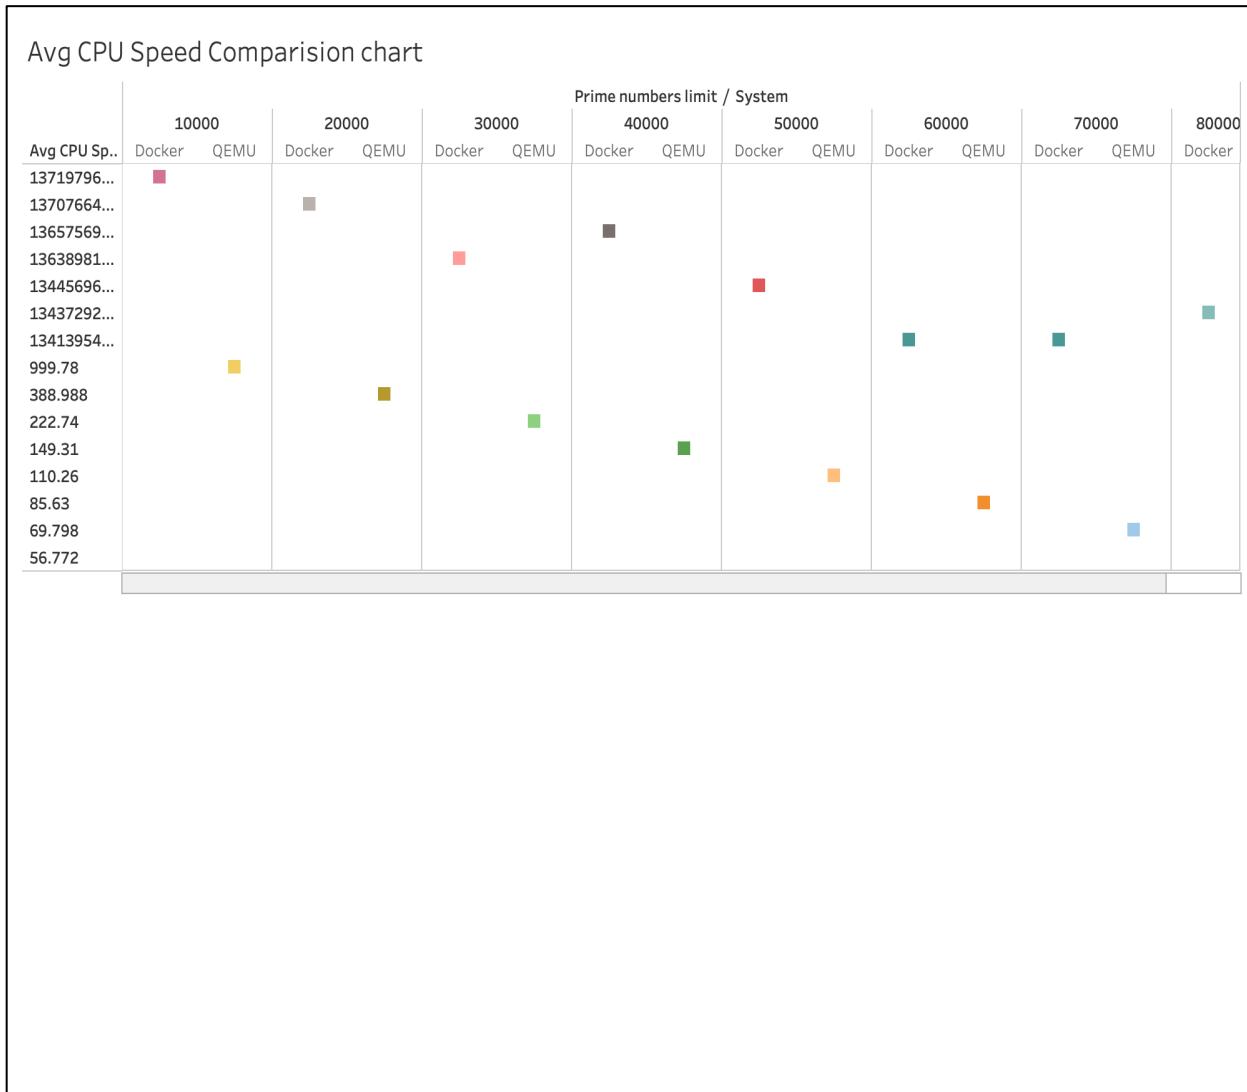
For Docker

Test run	CMP	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	80000	10.0078	11.96	17.54	22.76	56.93
2	80000	10.0126	8.5	17.7	27.83	56.41
3	80000	10.0035	17.35	17.62	21.19	56.66
4	80000	10.004	8.21	1738	20.41	57.46
5	80000	10.0147	11.53	17.71	23.45	56.4

For QEMU

Test run	CMP	Total time (In Sec)	Min (In ms)	Avg (in ms)	Max (in ms)	CPU Spped (Events /Sec)
1	80000	10.0078	11.96	17.54	22.76	56.93
2	80000	10.0126	8.5	17.7	27.83	56.41
3	80000	10.0035	17.35	17.62	21.19	56.66
4	80000	10.004	8.21	1738	20.41	57.46
5	80000	10.0147	11.53	17.71	23.45	56.4

CPU SPEED PERFORMANCE COMPARISION BETWEEN DOCKER AND QEMU VM



FILE IO PERFORMANCE TESTING:

- We are testing File IO performance using sysbench tool in both Ubuntu VM and Docker Container for below CPU parameters:

```
sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw prepare  
sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw run  
sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw cleanup
```

- Each Sysbench measurement is repeated at least 5 times for each test case

Automation for FILE IO performance test

- Automation for FILE IO performance test is done using Shell script program. And the logs are saved into a File named: FileIO_log_Docker.output .

```
1  #!/bin/bash  
2  #Loop  
3  exec 3>&1 4>&2  
4  trap 'exec 2>&4 1>&3' 0 1 2 3  
5  exec 1>FileIO_log_Docker.output 2>&1  
6  for j in {1..4}  
7  do  
8    for i in {1..4}  
9    do  
10   sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw prepare  
11   sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw run  
12   sysbench --num-threads=$j --test=fileio --file-total-size=3G --file-test-mode=rndrw cleanup  
13  
14  done  
15  done  
16
```

TEST CASE 1:

For Thread =1

```
sysbench --num-threads=1 --test=fileio --file-total-size=3G --file-test-mode=rndrw
prepare
sysbench --num-threads=1 --test=fileio --file-total-size=3G --file-test-mode=rndrw run
sysbench --num-threads=1 --test=fileio --file-total-size=3G --file-test-mode=rndrw
cleanup
```

Number of threads: 1 Initializing random number generator from current time
Extra file open flags: (none) 128 files, 24MiB each 36iB total file size Block size 16KiB Number of IO requests: 0 Read/Write ratio for combined random IO test: 1.50 Periodic FSYNC enabled, calling fsync() each 100 requests. Calling fsync() at the end of test, Enabled. Using synchronous I/O mode Doing random r/w test Initializing worker threads... Threads started!
File operations: reads/s: 10193.91 writes/s: 6795.91 fsyncs/s: 21758.75
Throughput: read, MiB/s: 159.28 written, MiB/s: 106.19
General statistics: total time: 10.0002s total number of events: 387380
Latency (ms): min: 0.00 avg: 0.03 max: 5.69 95th percentile: 0.00 sum: 9969.13
Threads fairness: events (avg/stddev): 387380.0000/0.00 execution time (avg/stddev): 9.9691/0.00

Number of threads: 1 Initializing random number generator from current time
Extra file open flags: (none) 128 files, 8MiB each 10iB total file size Block size 16KiB Number of IO requests: 0 Read/Write ratio for combined random IO test: 1.50 Periodic FSYNC enabled, calling fsync() each 100 requests. Calling fsync() at the end of test, Enabled. Using synchronous I/O mode Doing random r/w test Initializing worker threads... Threads started!
File operations: reads/s: 7141.04 writes/s: 4760.69 fsyncs/s: 15242.02
Throughput: read, MiB/s: 111.58 written, MiB/s: 74.39
General statistics: total time: 10.0063s total number of events: 271498
Latency (ms): min: 0.00 avg: 0.04 max: 3.13 95th percentile: 0.10 sum: 9931.79

Docker Container

Test run	Operation	FileSize	Threads	Throughput. - Read(MiB/Sec)	Throughput. written, MiB/s
1	r/w test	1G	1	159.28	106.19
2	r/w test	1G	1	141.64	94.42
3	r/w test	1G	1	137.83	91.89
4	r/w test	1G	1	134.22	89.48

QEMU VM

Test run	Operatio	FileSiz	Thread	Throughput. - Read(MiB/Sec)	Throughput. written, MiB
1	r/w test	1G	1	111.58	74.39
2	r/w test	1G	1	115.98	77.32
3	r/w test	1G	1	111.49	74.33
4	r/w test	1G	1	112.7	75.14
5	r/w test	1G	1	113.16	75.44

TEST CASE 2:

For Thread =4

```
sysbench --num-threads=4 --test=fileio --file-total-size=3G --file-test-mode=rndrw
prepare
sysbench --num-threads=4 --test=fileio --file-total-size=3G --file-test-mode=rndrw run
sysbench --num-threads=4 --test=fileio --file-total-size=3G --file-test-mode=rndrw
cleanup
```

<p>Running the test with following options: Number of threads: 4 Initializing random number generator from current time</p> <p>Extra file open flags: (none) 128 files, 24MiB each 3GiB total file size Block size 16KiB Number of IO requests: 0 Read/Write ratio for combined random IO test: 1.50 Periodic FSYNC enabled, calling fsync() each 100 requests. Calling fsync() at the end of test, Enabled. Using synchronous I/O mode Doing random r/w test Initializing worker threads...</p> <p>Threads started!</p> <p>File operations: reads/s: 22845.65 writes/s: 15230.23 fsyncs/s: 48784.10</p> <p>Throughput: read, MiB/s: 356.96 written, MiB/s: 237.97</p> <p>General statistics: total time: 10.0016s total number of events: 868257</p> <p>Latency (ms): min: 0.00 avg: 0.05 max: 9.56 95th percentile: 0.00 sum: 39915.74</p> <p>Threads fairness: events (avg/stddev): 217064.2500/509.32 execution time (avg/stddev): 9.9789/0.00</p>	<p>Number of threads: 4 Initializing random number generator from current time</p> <p>Extra file open flags: (none) 128 files, 8MiB each 16iB total file size Block size 16KiB Number of IO requests: 0 Read/Write ratio for combined random IO test: 1.50 Periodic FSYNC enabled, calling fsync() each 100 requests. Calling fsync() at the end of test, Enabled. Using synchronous I/O mode Doing random r/w test Initializing worker threads...</p> <p>Threads started!</p> <p>File operations: reads/s: 10571.50 writes/s: 7047.66 fsyncs/s: 22594.35</p> <p>Throughput: read, MiB/s: 165.18 written, MiB/s: 110.12</p> <p>General statistics: total time: 10.0170s total number of events: 402327</p> <p>Latency (ms): min: 0.00 avg: 0.10 max: 12.69 95th percentile: 0.26 sum: 39847.18</p>
--	---

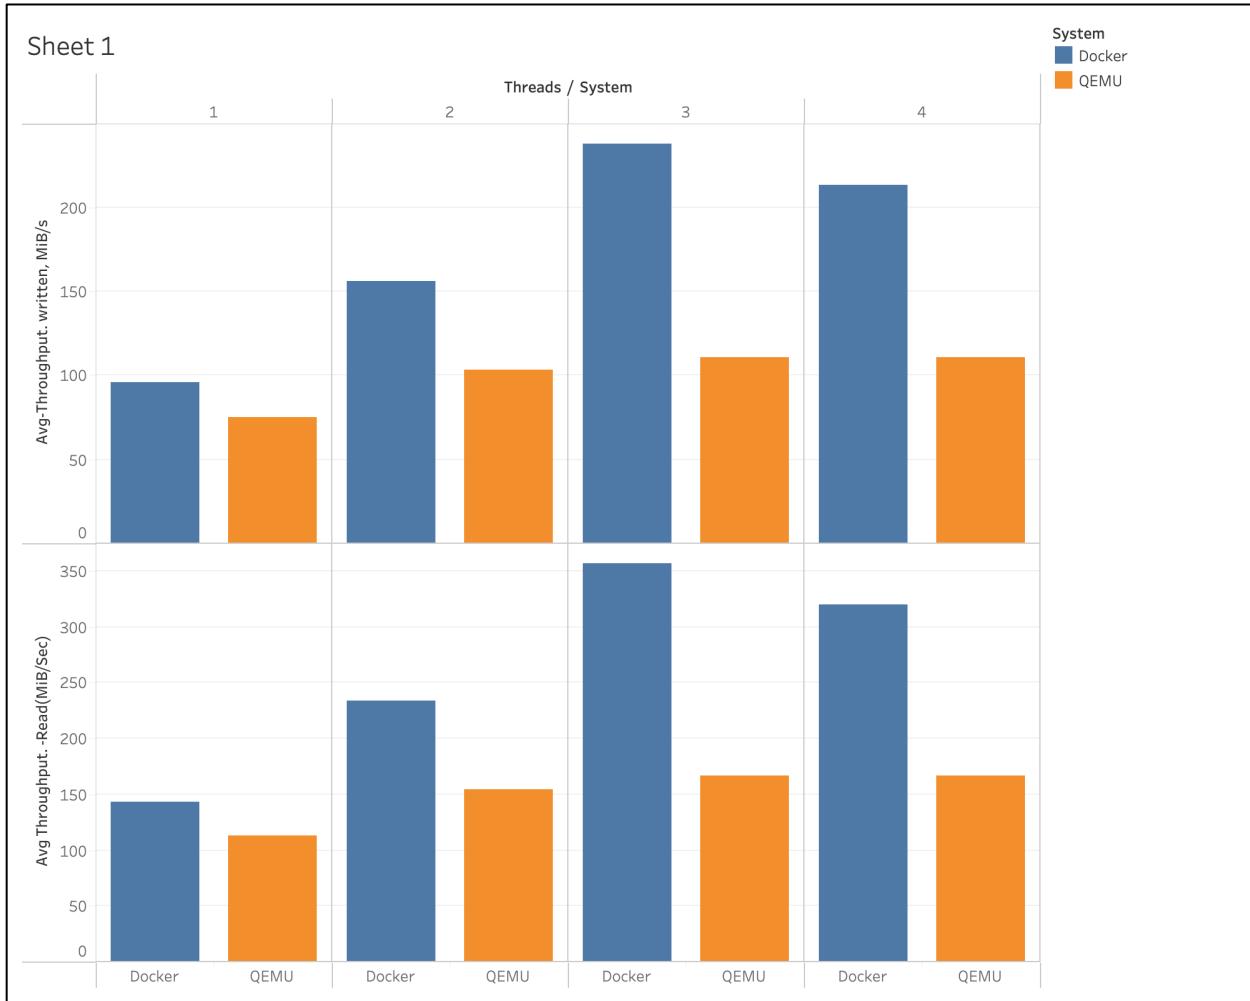
Docker container

Test run	Operation	FileSize	Threads	Throughput. - Read(MiB/Sec)	Throughput. written, MiB/s
1	r/w test	1G	4	166.43	110.95
2	r/w test	1G	4	166.22	110.82
3	r/w test	1G	4	166.44	110.96
4	r/w test	1G	4	166.44	110.95
5	r/w test	1G	4	165.18	110.12

QEMU VM

Test run	Operation	FileSize	Threads	Throughput. - Read(MiB/Sec)	Throughput. written, MiB/s
1	r/w test	1G	4	350.22	233.48
2	r/w test	1G	4	357.9	238.6
3	r/w test	1G	4	356.96	237.97
4	r/w test	1G	4	361.11	240.74

FILE I/O PERFORMANCE COMAPARISION B/W DOCKER AND QEMU



FILE I/O PERFORMANCE : AVERAGE THROUGHPUT

System	Threads	Avg Throughput. - Read(MiB/Sec)	Avg-Throughput. written, MiB/s
Docker	1	143.24	95.50
Docker	2	234.10	156.06
Docker	3	356.55	237.70
Docker	4	319.39	212.93
QEMU	1	112.982	75.324
QEMU	2	154.73	103.148
QEMU	3	166.338	110.892
QEMU	4	166.142	110.76

CONCLUSION:

1. The CPU performance was found to be MUCH better in Docker.
2. Avg CPU speed decrease with increase in Prime max limit. With Increase in threads performance increases.
3. FILE IO performance for docker is better than QEMU and with increase in thread performance improves
4. Average read and Written throughput increases with threads for both docker and QEMU.